

Prediction of individual sequences : homework

Lucas CLARTE

06 mars 2020

Part 1 - Link between online learning and game theory

All the code is written in the Python script `homework.py`.

1. For the game "Rock Paper scissors", we define the actions rock = 1, paper = 2, scissors = 3. The loss matrix L is as follows

$$L = \begin{pmatrix} 0 & 1 & -1 \\ -1 & 0 & 1 \\ 1 & -1 & 0 \end{pmatrix}$$

2. (a) We implement the function `rand_weighted(p)` that samples $i \in [M]$ with respect to the distribution $p \in \Delta_M$. In Python, the function is simply

```
def rand_weighted(p) :  
    return np.argmax(np.cumsum(p) > np.random.rand())[0]
```

In this function, we build an array `c = np.cumsum(p)` that contains the cumulative probabilities of p i.e $c_i = p_1 + \dots + p_i$. We then sample a number $x \in [0, 1)$ and determine the first index i in the cumulative array `c` such that $c_i > x$. It is easy to show that for all j , $i = j$ with probability p_j .

2. (b) To implement the function `EWA_update(p, l)` we simply multiply the vectors p and $\exp(-\eta l(i))$ component-wise and then normalize the new vector.

3. (a.) We simulate EWA against a fixed adversary with strategy $q = (0.5, 0.25, 0.25)$, in the function `EWA(L, T, eta, q)`. This function runs T iterations. At iteration t , we sample an action j_t from q using `rand_weighted` and update the weights p_t with the function `EWA_update` called with the loss l_t . The loss $l_t(i)$ of the player if he choses the action i is equal to $L(i, j_t)$.

3. (b) As asked, we simulate the game with $T = 100, \eta = 1.0$ and plot the weight vectors p_1, \dots, p_T in the figure (??) . We see that the best strategy is $p = (0.0, 1.0, 0.0)$. We can prove it rigorously. Indeed, consider the strategy $p = (x, y, z)$ that minimizes the average loss $l(p, q) = \mathbb{E}_{i \sim p, j \sim q}(L(i, j))$. We write

$$\begin{aligned} l(p, q) &= x \times (0.25 - 0.25) + y \times (0.25 - 0.5) + z \times (0.5 - 0.25) \\ &= -0.25 \times y + 0.25 \times z \end{aligned}$$

The optimal value of p that satisfies the constraint $p \in \Delta_3$ and minimizes the above expression is $x = 0, y = 1, z = 0$ which is what we wanted.

3. (c) We plot the average loss $\bar{l}_t = \frac{1}{t} \sum_{1 \leq s \leq t} l(i_s, j_s)$ as a function of t , and obtain the figure (??). The figure shows 10 different runs of the simulation.

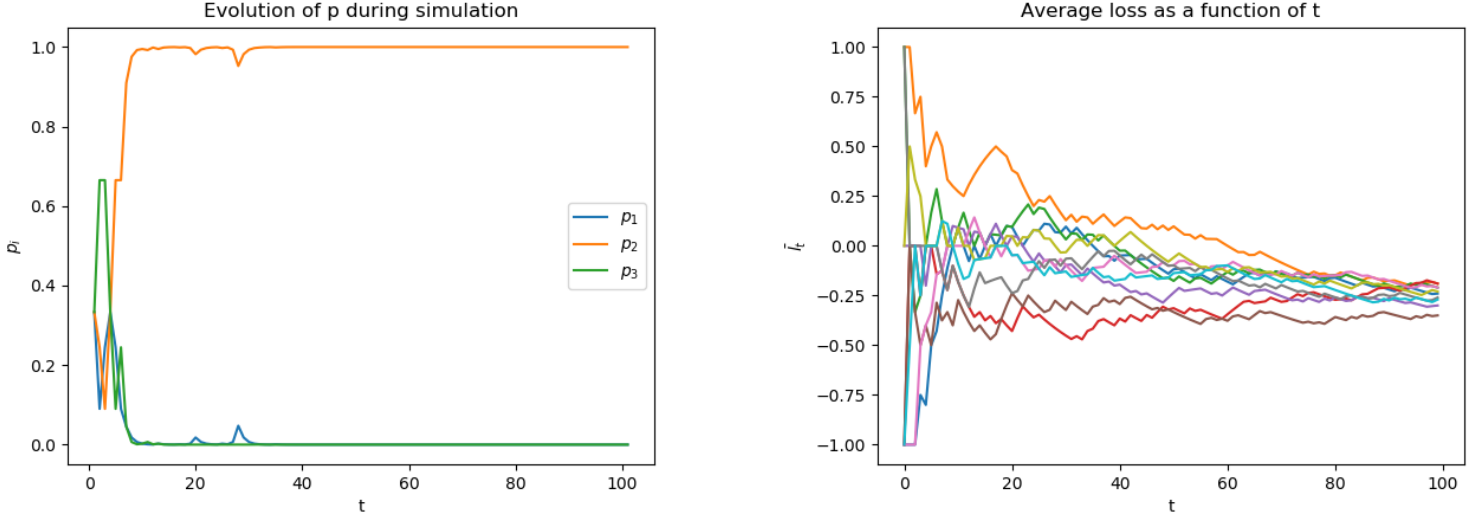


Figure 1: Left : plot of the player's strategy p with EWA, when opponent has a fixed strategy $(0.5, 0.25, 0.25)$. Right : plot of the average loss as a function of t for 10 simulations of the game with EWA

3. (d) Here, the cumulative regret is defined by

$$R_t = \sum_{s=1}^t L(i_s, j_s) - \min_i \sum_{s \leq t} L(i, j_s) \quad (1)$$

Figure (??) shows the cumulative regret for 10 trials with $T = 100$. We observe experimentally that for each trial, after a certain time the cumulative regret is constant.

3. (e) To estimate the stability, $n = 10$ simulations are executed, and we plot in figure (??) the n average losses as well as the minimum, maximum and mean losses. At first glance, the EWA seems to be stable. Indeed, the maximum and minimum both seem to converge towards the value $l_\infty = -0.25$.

3. (f) In theory, the best learning rate η with EWA is $\eta_{\text{EWA}} = \sqrt{2\ln(K)/T}$ with $K = 3$ here and $T = 100$, so we obtain $\eta_{\text{EWA}} \simeq 0.15$. We see that in practice, the best learning rate is not necessarily equal to η_{EWA} .

4. Simulation against an adaptive adversary In this question, the player still uses EWA with a learning rate $\eta = 1.0$ while the opponent uses a learning rate $\eta = 0.05$

4.(a) When the adversary uses EWA like the player, we observe that the average loss seems to converge towards 0 which is the value of the game.

4. (b)




Figure 2: Left : Cumulative regret .Right : average, minimum and maximum of the average loss \bar{l}_t for $n = 10$ simulation trials. The n average losses are plotted in grey.

Bandit feedback We now assume that the player doesn't have access to L but only the incurred loss $L(i_t, j_t)$. We will here implement the algorithm **Exp3** where at each time t , the action i_t is selected with a probability

$$\mathbb{P}(i_t = i) = \frac{\exp(-\eta \hat{l}_t(i))}{\sum_j \exp(-\eta \hat{l}_t(j))} \quad (2)$$

where $\hat{l}_t(i)$ is the *estimated loss* of the action i , i.e the average loss incurred by the player when the action i was selected.

5.(a) To implement the function **estimated_loss**, we maintain a vector $C \in \mathbb{N}^M$ where at each time, C_i is the number of times the action i has been played. We also maintain the vector $S \in \mathbb{R}^M$ where $S_i = \sum_j L(i_t, j_t) \mathbf{1}(i_t = i)$ is the sum of the losses occurred by the player every time they chose the action i . The function **estimated_loss(i)** simply returns S_i/C_i .

5. (b)




Figure 3: Average loss for $n = 10$ simulations with adaptive adversary

